

Hungryserv

A Homeserver Optimized for Unfederated Use-Cases

Sumner Evans

26 August 2022

Beeper

A bit about me

My name is Sumner, I'm a **software engineer at Beeper**.

- I graduated from Colorado School of Mines in 2019 with a master's in CS.
- I teach as an adjunct professor at my alma mater.
- I enjoy skiing, volleyball, and football (soccer).
- I'm a 4th degree black belt in ATA taekwondo.

I became interested in Matrix when I was the chair of the ACM chapter at Mines. I was looking for an open source chat platform for the club to use, and Matrix fit the bill!

A bit about me

My name is Sumner, I'm a **software engineer at Beeper**.

- I graduated from Colorado School of Mines in 2019 with a master's in CS.
- I teach as an adjunct professor at my alma mater.
- I enjoy skiing, volleyball, and football (soccer).
- I'm a 4th degree black belt in ATA taekwondo.

I became interested in Matrix when I was the chair of the ACM chapter at Mines. I was looking for an open source chat platform for the club to use, and Matrix fit the bill!

What I work on at Beeper

I am on the newly created *Scaling* team.

Our current objective is to prepare Beeper for rocket-ship growth.

I was previously part of the *Bridges* team.

Notable projects included:

- Writing the LinkedIn bridge
- Adding bridge status reporting and message send status.
- Implementing massive stability improvements in the Signal bridge
- Implementing incremental infinite backfill in our WhatsApp and Facebook bridges

What I work on at Beeper

I am on the newly created *Scaling* team.

Our current objective is to prepare Beeper for rocket-ship growth.

I was previously part of the *Bridges* team.

Notable projects included:

- Writing the LinkedIn bridge
- Adding bridge status reporting and message send status.
- Implementing massive stability improvements in the Signal bridge
- Implementing incremental infinite backfill in our WhatsApp and Facebook bridges

Overview

1. A bit about Beeper
2. A bit about Beeper's current architecture
3. Our new architecture

This talk is interactive!

If you have questions at any point, feel free to interrupt me.

A bit about Beeper

Our mission is to:

make it easy for everyone on Earth to chat with each other.

We specifically chose the word “chat” rather than “communicate” because we are focusing on *people talking to one another*.


Our mission is to:

make it easy for everyone on Earth to chat with each other.

We specifically chose the word “chat” rather than “communicate” because we are focusing on *people talking to one another*.

What is Beeper?

Beeper is an app that brings all of your chat networks together into a single inbox.



What is Beeper?

Beeper allows you to consolidate messages from 15+ chat apps into a single inbox.

- Beeper is available on macOS, Windows, Linux iPhone, iPad, Android and Chrome OS.
- Beeper is built on top of the Matrix protocol.
- Beeper encrypts all chats, including bridged chats, by default¹.

¹We have to momentarily decrypt your chat messages to translate from the other network to Matrix, but we never log or store your unencrypted messages.

What is Beeper?

Beeper allows you to consolidate messages from 15+ chat apps into a single inbox.

- Beeper is available on macOS, Windows, Linux iPhone, iPad, Android and Chrome OS.
- Beeper is built on top of the Matrix protocol.
- Beeper encrypts all chats, including bridged chats, by default¹.

¹We have to momentarily decrypt your chat messages to translate from the other network to Matrix, but we never log or store your unencrypted messages.

Bridging to a glorious Matrix future

We see bridges as a way to onboard users into the Matrix ecosystem.

Users can switch to a new app without losing a single conversation!

We encourage our users to connect *all* of their chat networks to Beeper.

Bridging to a glorious Matrix future

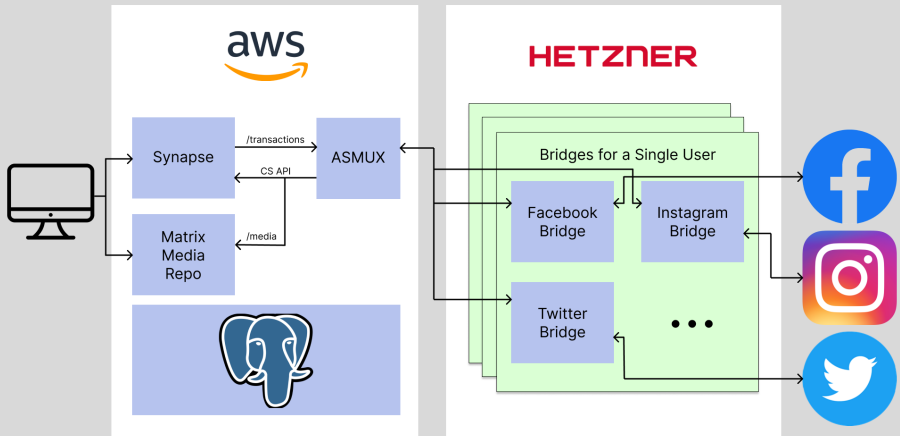
We see bridges as a way to onboard users into the Matrix ecosystem.

Users can switch to a new app without losing a single conversation!

We encourage our users to connect *all* of their chat networks to Beeper.

A bit about Beeper's current architecture

A diagram



Each user gets their own bridge for each network they connect.

Advantages of our architecture

- ASMUX allows us to dynamically change what bridges are running
- Each users' bridge is **isolated**
- We can deploy **different bridge versions** to different sets of users
- Bridges can be stopped and started individually
- Double puppeting

Advantages of our architecture

- ASMUX allows us to dynamically change what bridges are running
- Each users' bridge is **isolated**
- We can deploy **different bridge versions** to different sets of users
- Bridges can be stopped and started individually
- Double puppeting

Advantages of our architecture

- ASMUX allows us to dynamically change what bridges are running
- Each users' bridge is **isolated**
- We can deploy **different bridge versions** to different sets of users
- Bridges can be stopped and started individually
- Double puppeting

Advantages of our architecture

- ASMUX allows us to dynamically change what bridges are running
- Each users' bridge is **isolated**
- We can deploy **different bridge versions** to different sets of users
- Bridges can be stopped and started individually
- Double puppeting

Advantages of our architecture

- ASMUX allows us to dynamically change what bridges are running
- Each users' bridge is **isolated**
- We can deploy **different bridge versions** to different sets of users
- Bridges can be stopped and started individually
- Double puppeting

Disadvantages of our architecture

- We have to run a **lot** of bridges
- Synapse is not designed for a dynamic numbers of bridges
- If two users join the same chat on an external network, we end up with two rooms with the same data
- We run Synapse on AWS which is relatively expensive
- Synapse becomes a bottleneck for all traffic

Disadvantages of our architecture

- We have to run a **lot** of bridges
- Synapse is not designed for a dynamic numbers of bridges
- If two users join the same chat on an external network, we end up with two rooms with the same data
- We run Synapse on AWS which is relatively expensive
- Synapse becomes a bottleneck for all traffic

Disadvantages of our architecture

- We have to run a **lot** of bridges
- Synapse is not designed for a dynamic numbers of bridges
- If two users join the same chat on an external network, we end up with two rooms with the same data
- We run Synapse on AWS which is relatively expensive
- Synapse becomes a bottleneck for all traffic


Disadvantages of our architecture

- We have to run a **lot** of bridges
- Synapse is not designed for a dynamic numbers of bridges
- If two users join the same chat on an external network, we end up with two rooms with the same data
- We run Synapse on AWS which is relatively expensive
- Synapse becomes a bottleneck for all traffic

Disadvantages of our architecture

- We have to run a **lot** of bridges
- Synapse is not designed for a dynamic numbers of bridges
- If two users join the same chat on an external network, we end up with two rooms with the same data
- We run Synapse on AWS which is relatively expensive
- Synapse becomes a bottleneck for all traffic

Back to the diagram



The core software in this architecture is Synapse.

Because we encourage users to connect every chat network, we have a lot of puppet users.

On average, each user brings 4716 puppeted users.

All of these puppets represent real people on other networks, so they each generate their own traffic.

On average, each user and their puppets collectively account for 100k events!

None of that traffic is federated, but it has to go to Synapse which is designed with federation front-and-center!

Because we encourage users to connect every chat network, we have a lot of puppet users.

On average, each user brings 4716 puppeted users.

All of these puppets represent real people on other networks, so they each generate their own traffic.

On average, each user and their puppets collectively account for 100k events!

None of that traffic is federated, but it has to go to Synapse which is designed with federation front-and-center!

Because we encourage users to connect every chat network, we have a lot of puppet users.

On average, each user brings 4716 puppeted users.

All of these puppets represent real people on other networks, so they each generate their own traffic.

On average, each user and their puppets collectively account for 100k events!

None of that traffic is federated, but it has to go to Synapse which is designed with federation front-and-center!

A few more notes

- **Synapse is aggressively federated**

Synapse has no special handling of federated vs unfederated rooms.

- **Synapse does not handle message floods well**

Synapse falls over when users bridge some large Telegram chats and Discord servers.

- **Deleting rooms from Synapse is hard**

We end up with lots of *dead* rooms when users delete their bridges and never turn it back on (or need a hard bridge reset).

A few more notes

- **Synapse is aggressively federated**

Synapse has no special handling of federated vs unfederated rooms.

- **Synapse does not handle message floods well**

Synapse falls over when users bridge some large Telegram chats and Discord servers.

- **Deleting rooms from Synapse is hard**

We end up with lots of *dead* rooms when users delete their bridges and never turn it back on (or need a hard bridge reset).

- **Synapse is aggressively federated**

Synapse has no special handling of federated vs unfederated rooms.

- **Synapse does not handle message floods well**

Synapse falls over when users bridge some large Telegram chats and Discord servers.

- **Deleting rooms from Synapse is hard**


We end up with lots of *dead* rooms when users delete their bridges and never turn it back on (or need a hard bridge reset).

Solution: shard our homeserver based on
bridged vs. unbridged traffic

Our new architecture



A sharded future



Let's talk about the differences from the old architecture...

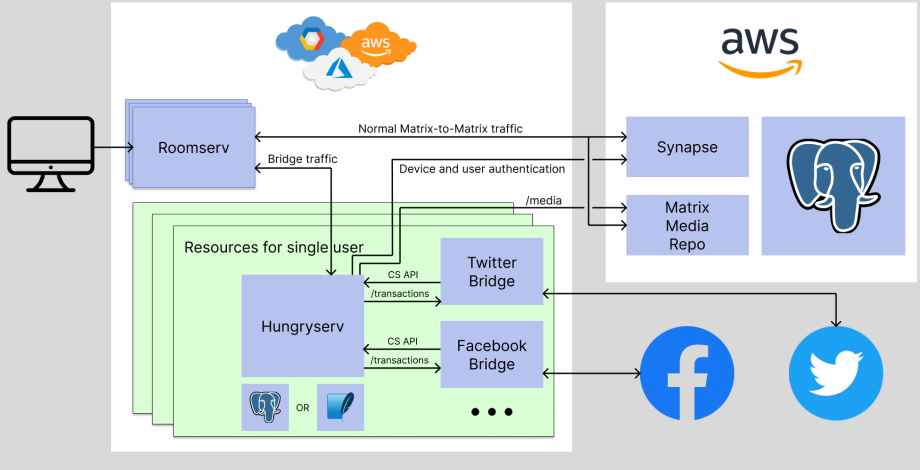
This new architecture allows us to handle Matrix-native federated traffic separately from our high-throughput unfederated bridge traffic.

All while maintaining API compatibility with clients, bridges, and the Matrix federation!

This new architecture allows us to handle Matrix-native federated traffic separately from our high-throughput unfederated bridge traffic.

All while maintaining API compatibility with clients, bridges, and the Matrix federation!

A sharded future



Let's talk about the differences from the old architecture...

Hungryserv: why is it hungry?

Because it's unfederated!

Hungryserv: why is it hungry?

Because it's unfederated!

We run a **Hungryserv** instance for every user to handle that user's bridge traffic.

- If there's a message flood from a bridge, it only affects that user.
- We don't have to store all of the bridge events and users in our Synapse database.
- Synapse is less resource constrained and can do what it does best: federate with other Matrix homeservers.

We are also currently looking into ways of using **SQLite** in production using Litestream.

We run a **Hungryserv** instance for every user to handle that user's bridge traffic.

- If there's a message flood from a bridge, it only affects that user.
- We don't have to store all of the bridge events and users in our Synapse database.
- Synapse is less resource constrained and can do what it does best: federate with other Matrix homeservers.

We are also currently looking into ways of using **SQLite** in production using Litestream.

We run a **Hungryserv** instance for every user to handle that user's bridge traffic.

- If there's a message flood from a bridge, it only affects that user.
- We don't have to store all of the bridge events and users in our Synapse database.
- Synapse is less resource constrained and can do what it does best: federate with other Matrix homeservers.

We are also currently looking into ways of using **SQLite** in production using Litestream.

We run a **Hungryserv** instance for every user to handle that user's bridge traffic.

- If there's a message flood from a bridge, it only affects that user.
- We don't have to store all of the bridge events and users in our Synapse database.
- Synapse is less resource constrained and can do what it does best: federate with other Matrix homeservers.

We are also currently looking into ways of using **SQLite** in production using Litestream.

We run a **Hungryserv** instance for every user to handle that user's bridge traffic.

- If there's a message flood from a bridge, it only affects that user.
- We don't have to store all of the bridge events and users in our Synapse database.
- Synapse is less resource constrained and can do what it does best: federate with other Matrix homeservers.

We are also currently looking into ways of using **SQLite** in production using Litestream.

Hungryserv: aggressively unfederated

Each bridge belongs to a single user, so federation is unnecessary.

When your DAG is a linked-list, don't store it as a DAG.

Being unfederated means that:

- we can avoid implementing the state resolution algorithm,
- event storage requires less metadata,
- deletions can be optimized, and
- infinite backfill becomes an simple batch SQL operation.

Hungryserv: aggressively unfederated

Each bridge belongs to a single user, so federation is unnecessary.

When your DAG is a linked-list, don't store it as a DAG.

Being unfederated means that:

- we can avoid implementing the state resolution algorithm,
- event storage requires less metadata,
- deletions can be optimized, and
- infinite backfill becomes an simple batch SQL operation.

Hungryserv: aggressively unfederated

Each bridge belongs to a single user, so federation is unnecessary.

When your DAG is a linked-list, don't store it as a DAG.

Being unfederated means that:

- we can avoid implementing the state resolution algorithm,
- event storage requires less metadata,
- deletions can be optimized, and
- infinite backfill becomes an simple batch SQL operation.

Goals

- Be fast
- Be compatible with Beeper clients (and mostly compatible with Element clients)
- Be API-compatible with the Appservice API for bridges

Non-goals

- Be federated
- Be fully spec compliant


Goals

- Be fast
- Be compatible with Beeper clients (and mostly compatible with Element clients)
- Be API-compatible with the Appservice API for bridges

Non-goals

- Be federated
- Be fully spec compliant

A sharded future




Let's talk about the differences from the old architecture...

Roomserv is a service which splits client requests between Synapse and Hungryserv.

Roomserv merges the responses from the two homeservers to make it appear as a single homeserver to the client.

A sharded future



Let's talk about the differences from the old architecture...

No more ASMUX

Hungryserv natively supports dynamic registration of app services, and knows how to route requests to the correct bridge.

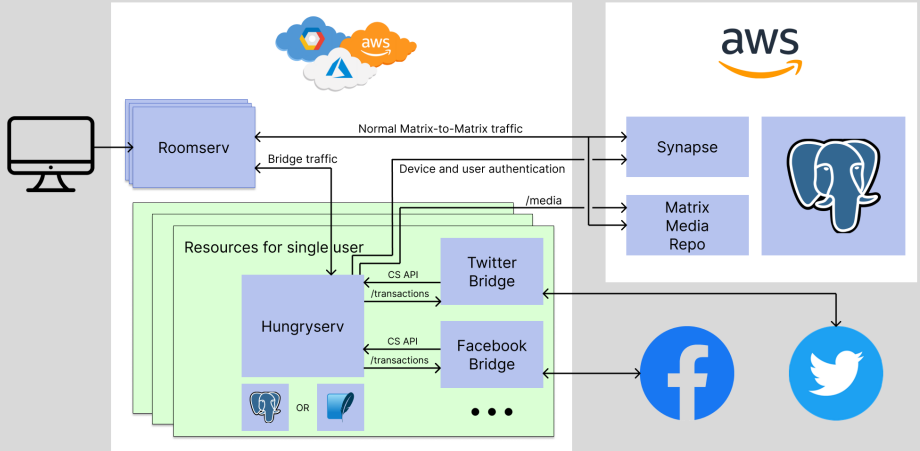
We no longer have to rely on the single transactions firehose from Synapse.

No more ASMUX

Hungryserv natively supports dynamic registration of app services, and knows how to route requests to the correct bridge.

We no longer have to rely on the single transactions firehose from Synapse.

A sharded future



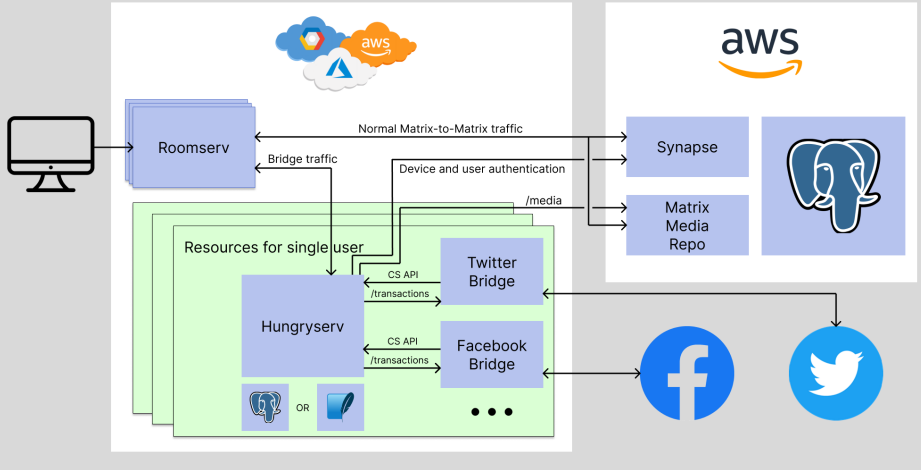
Let's talk about the differences from the old architecture...

A multi-cloud future?

The resources for each user can be moved closer to them using a multi-cloud, or multi-region strategy.

This will help reduce latency (especially to bridged networks) and allow for us to deploy in more flexible and cost-effective environments.

A sharded future



Let's talk about the differences from the old architecture...

Demo!

Questions?